

Algorithms

Chapter 13 Red-Black Trees

Associate Professor: Ching-Chi Lin

林清池 副教授

chingchi.lin@gmail.com

Department of Computer Science and Engineering
National Taiwan Ocean University

Outline


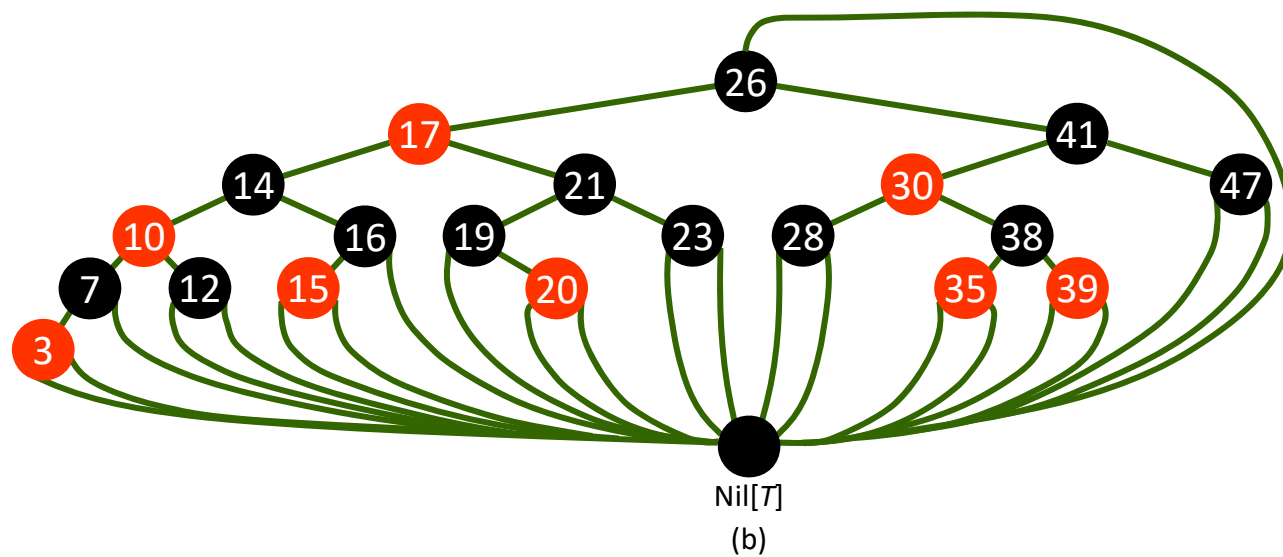

- ▶ **Properties of red-black trees**
- ▶ Rotations
- ▶ Insertion
- ▶ Deletion

Overview

- ▶ A binary search tree of height h can implement any of the basic dynamic-set operations such as SEARCH, PREDECESSOR, SUCCESSOR, MINIMUM, MAXIMUM, INSERT, and DELETE in $O(h)$ time.
- ▶ Thus, the set operations are fast if the height of the search tree is small; but if its height is large, their performance may be no better than with a linked list.
- ▶ **Red-black trees**
 - ▶ A variation of binary search trees.
 - ▶ **Balanced**: height is $O(\lg n)$, where n is the number of nodes.
 - ▶ Operations will take $O(\lg n)$ time in the worst case.

Properties of red-black trees_{1/2}

- ▶ A **red-black tree** = a binary search tree + 1 bit per node: an attribute **color**, which is either red or black.
 - ▶ Each node of the tree now contains the fields **color**, **key**, **left**, **right**, and **p**.
 - ▶ If a child or the parent of a node does not exist, the corresponding pointer field of the node contains the value NIL.
 - ▶ **Red-black properties**
 1. Every node is either red or black.
 2. The root is black.
 3. Every leaf (Nil) is black.
 4. If a node is red, then both its children are black.
 5. For each node, all paths from the node to descendant leaves contain the same number of black nodes.
-



Properties of red-black trees_{2/2}

- ▶ All leaves are empty (nil) and colored black.
 - ▶ We use a single sentinel, $\text{nil}[T]$, for all the leaves as a matter of convenience.
 - ▶ $\text{color}[\text{nil}[T]]$ is black.
 - ▶ The root's parent is also $\text{nil}[T]$.
- ▶ **Height of a red-black tree**
 - ▶ **Height of a node** is the number of edges in a longest path to a leaf.
 - ▶ **Black-height** of a node x : $\text{bh}(x)$ is the number of black nodes (including $\text{nil}[T]$) on the path from x to leaf, not counting x . By property 5, black-height is well defined.

The height of a red-black tree_{1/2}

- ▶ **Claim 1** The subtree rooted at any node x contains $\geq 2^{\text{bh}(x)} - 1$ internal nodes.

Proof: By induction on the height of x .

- ▶ **Basis:**

- ▶ Height of $x = 0 \Rightarrow x$ is a leaf $\Rightarrow \text{bh}(x) = 0$.
- ▶ The subtree rooted at x has 0 internal nodes. $2^0 - 1 = 0$.

- ▶ **Inductive step:**

- ▶ Let $\text{bh}(x) = b$.
- ▶ Any child of x has black-height either b (if the child is red) or $b - 1$ (if the child is black).
- ▶ By the inductive hypothesis, each child has $\geq 2^{\text{bh}(x)-1} - 1$ internal nodes.
- ▶ Thus, the subtree rooted at x contains $\geq 2 \cdot (2^{\text{bh}(x)-1} - 1) + 1 = 2^{\text{bh}(x)} - 1$ internal nodes. (The +1 is for x itself.)

The height of a red-black tree_{2/2}

- ▶ **Lemma 1** A red-black tree with n internal nodes has height $\leq 2\lg(n+1)$.

Proof:

- ▶ Let h and b be the height and black-height of the root, respectively.
- ▶ By Claim 1, we have $n \geq 2^b - 1$.
- ▶ By property 4, $\leq h/2$ nodes on the path from the node to a leaf are red.
- ▶ Hence $\geq h/2$ are black, i.e., $b \geq h/2$.
- ▶ Thus, $n \geq 2^b - 1 \geq 2^{h/2} - 1$.
- ▶ This implies $h \leq 2\lg(n+1)$.

Operations on red-black trees


- ▶ The non-modifying binary-search-tree operations MINIMUM, MAXIMUM, SUCCESSOR, PREDECESSOR, and SEARCH run in $O(\text{height})$ time. Thus, they take $O(\lg n)$ time on red-black trees.
- ▶ **Insertion** and **deletion** are not so easy.
- ▶ For example:
 - ▶ If we insert, what color to make the new node?
 - ▶ Red? Might violate property 4.
 - ▶ Black? Might violate property 5.

Outline

- ▶ Properties of red-black trees
- ▶ **Rotations**
- ▶ Insertion
- ▶ Deletion

Rotations

- ▶ A local operation in a search tree that preserves the binary-search-tree property.
- ▶ There are two kinds of rotations:
 - ▶ **Left rotations** and **right rotations**.
 - ▶ They are inverses of each other.
- ▶ When we do a left rotation on a node x , we assume that its right child y is not $\text{nil}[T]$.




LEFT-ROTATE pseudocode

LEFT-ROTATE(T, x)

1. $y \leftarrow \text{right}[x]$
2. $\text{right}[x] \leftarrow \text{left}[y]$
3. **if** $\text{left}[y] \neq \text{nil}[T]$
4. $p[\text{left}[y]] \leftarrow x$
5. $p[y] \leftarrow p[x]$
6. **if** $p[x] = \text{nil}[T]$
7. $\text{root}[T] \leftarrow y$
8. **else if** $x = \text{left}[p[x]]$
9. $\text{left}[p[x]] \leftarrow y$
10. **else** $\text{right}[p[x]] \leftarrow y$
11. $\text{left}[y] \leftarrow x$
12. $p[x] \leftarrow y$

- ▶ Time: $O(1)$.
- ▶ The code for RIGHT-ROTATE is symmetric.



Outline


- ▶ Properties of red-black trees
- ▶ Rotations
- ▶ **Insertion**
- ▶ Deletion

RB-Insertion_{1/2}

- ▶ Start by doing regular binary-search-tree insertion.

RB-INSERT(T, z)

1. $y \leftarrow \text{NIL}; x \leftarrow \text{root}[T]$
2. **while** $x \neq \text{NIL}$
3. $y \leftarrow x$
4. **if** $\text{key}[z] < \text{key}[x]$
5. $x \leftarrow \text{left}[x]$
6. **else** $x \leftarrow \text{right}[x]$
7. $p[z] \leftarrow y$
8. **if** $y = \text{NIL}$
9. $\text{root}[T] \leftarrow z$ /* Tree T was empty */
10. **else if** $\text{key}[z] < \text{key}[y]$
11. $\text{left}[y] \leftarrow z$
12. **else** $\text{right}[y] \leftarrow z$
13. $\text{left}[z] \leftarrow \text{nil}[T]; \text{right}[z] \leftarrow \text{nil}[T]; \text{color}[z] \leftarrow \text{RED}$
14. RB-INSERT-FIXUP(T, z)



RB-Insertion_{2/2}


- ▶ RB-INSERT ends by coloring the new node z red.
- ▶ Then it calls RB-INSERT-FIXUP because we could have violated a red-black property.
- ▶ Which property might be violated?
 1. OK.
 2. **If z is the root**, then there's a violation. Otherwise, OK.
 3. OK.
 4. **If $p[z]$ is red**, there's a violation: both z and $p[z]$ are red.
 5. OK.

RB-INSERT-FIXUP procedure


RB-INSERT-FIXUP(T, z)

1. **while** $color[p[z]] = \text{RED}$
 2. **if** $p[z] = \text{left}[p[p[z]]]$
 3. $y \leftarrow \text{right}[p[p[z]]]$
 4. **if** $color[y] = \text{RED}$
 5. $color[p[z]] \leftarrow \text{BLACK}$ Case 1
 6. $color[y] \leftarrow \text{BLACK}$ Case 1
 7. $color[p[p[z]]] \leftarrow \text{RED}$ Case 1
 8. $z \leftarrow p[p[z]]$ Case 1
 9. **else** {
 10. **if** $z = \text{right}[p[z]]$ Case 2
 11. $z \leftarrow p[z]$ Case 2
 12. LEFT-ROTATE(T, z) Case 2
 13. }
 14. $color[p[z]] \leftarrow \text{BLACK}$ Case 3
 15. $color[p[p[z]]] \leftarrow \text{RED}$ Case 3
 16. RIGHT-ROTATE($T, p[p[z]]$) Case 3
 17. **else** (same as **then** clause
 18. with “right” and “left” exchanged)
 19. $color[\text{root}[T]] \leftarrow \text{BLACK}$
-


(a)




(b)



(c)



(d)



Correctness of RB-INSERT

- ▶ **Loop invariant:** At the start of each iteration of the **while** loop of lines 1-16,
 - a. Node z is red.
 - b. If $p[z]$ is the root, then $p[z]$ is black.
 - c. There is at most one red-black violation:
 - ▶ Property 2, z is the root and is red.
 - ▶ Property 4, both z and $p[z]$ are red.

- ▶ We omit the further details for proving the correctness.

Time complexity of RB-INSERT


▶ **Analysis:**

- ▶ Each iteration takes $O(1)$ time.
 - ▶ The **while** loop repeats only if case 1 is executed, and then the pointer z moves two levels up the tree.
 - ▶ The **while** loop terminates if case 2 or case 3 is executed.
 - ▶ $O(\lg n)$ levels $\rightarrow O(\lg n)$ time.
 - ▶ Also note that there are at most 2 rotations overall.
-
- ▶ Thus, insertion into a red-black tree takes $O(\lg n)$ time.

Outline

- ▶ Properties of red-black trees
- ▶ Rotations
- ▶ Insertion
- ▶ **Deletion**


Case 1:



Case 2:



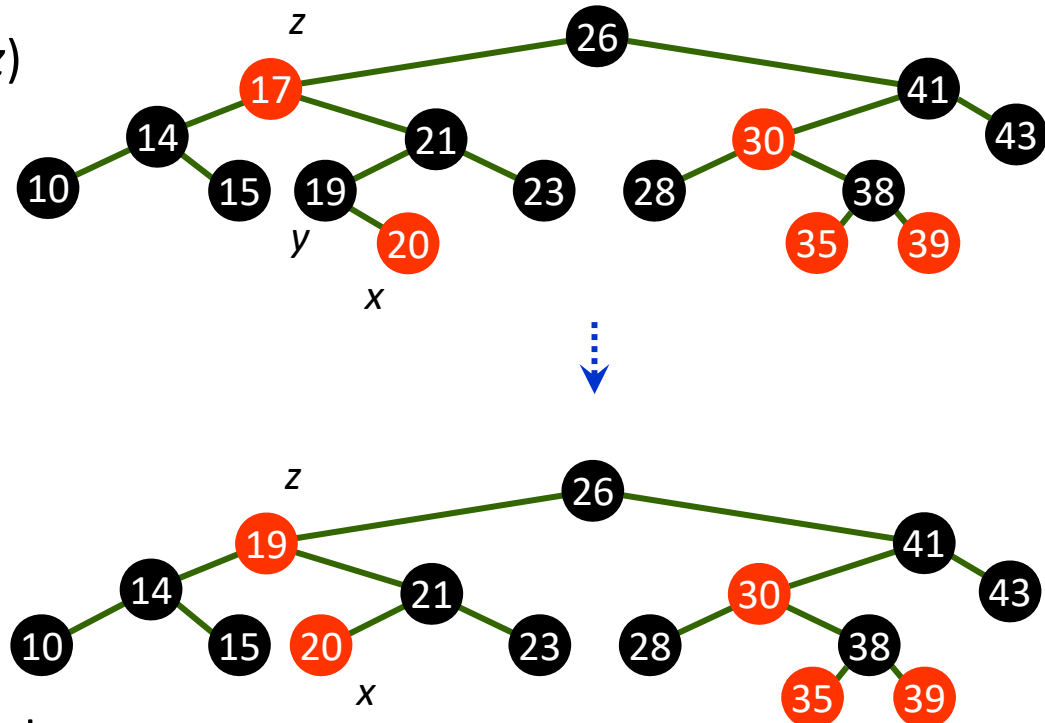
Case 3:



RB-Deletion_{1/3}

RB-DELETE(T, z)

1. **if** $left[z] = \text{NIL}$ or $right[z] = \text{NIL}$
2. $y \leftarrow z$
3. **else** $y \leftarrow \text{TREE-SUCCESSOR}(z)$
4. **if** $left[y] \neq \text{NIL}$
5. $x \leftarrow left[y]$
6. **else** $x \leftarrow right[y]$
7. $p[x] \leftarrow p[y]$
8. **if** $p[y] = \text{NIL}$
9. $root[T] \leftarrow x$
10. **else if** $y = left[p[y]]$
11. $left[p[y]] \leftarrow x$
12. **else** $right[p[y]] \leftarrow x$
13. **if** $y \neq z$
14. $key[z] \leftarrow key[y]$
15. copy y 's satellite data into z
16. **if** $color[y] = \text{BLACK}$
17. RB-DELETE-FIXUP(T, x)
18. **return** y



RB-Deletion_{2/3}

- ▶ y is the node that was actually spliced out.
- ▶ x is either
 - ▶ y 's sole non-sentinel child before y was spliced out, or
 - ▶ the sentinel, if y had no children.
- ▶ In both cases, $p[x]$ is now the node that was previously y 's parent.
- ▶ If y is red, the red-black properties still hold when y is spliced out, for the following reasons:
 - ▶ no black-heights in the tree have changed,
 - ▶ no red nodes have been made adjacent, and
 - ▶ since y could not have been the root if it was red, the root remains black.

RB-Deletion_{3/3}

- ▶ If y is black, we could have violations of red-black properties:
 1. OK.
 2. If y is the root and x is red, then the root has become red.
 3. OK.
 4. Violation if $p[y]$ and x are both red.
 5. Any path containing y now has 1 fewer black node.
- ▶ Correct this problem by giving x an “extra black”.
 - ▶ Now property 5 is OK, but property 1 is not.
 - ▶ x is either **doubly black** or **red & black**.

RB-DELETE-FIXUP


- ▶ **Idea:** Move the extra black up the tree until
 1. x points to a red & black node \rightarrow turn it into a black node,
 2. x points to the root \rightarrow just remove the extra black, or
 3. suitable rotations and recolorings can be performed.
- ▶ Within the while loop:
 - ▶ x always points to a nonroot doubly black node.
 - ▶ w is x 's sibling.
 - ▶ w cannot be $nil[T]$, since that would violate property 5 at $p[x]$.
- ▶ There are 8 cases, 4 of which are symmetric to the other 4.
- ▶ As with insertion, the cases are not mutually exclusive.
We'll look at cases in which x is a left child.

RB-DELETE-FIXUP procedure

RB-DELETE-FIXUP(T, x)


```
1.  while  $x \neq \text{root}[T]$  and  $\text{color}[x] = \text{BLACK}$ 
2.      if  $x = \text{left}[p[x]]$ 
3.           $w \leftarrow \text{right}[p[x]]$ 
4.          if  $\text{color}[w] = \text{RED}$ 
5.               $\text{color}[w] \leftarrow \text{BLACK}$  Case 1
6.               $\text{color}[p[x]] \leftarrow \text{RED}$  Case 1
7.              LEFT-ROTATE( $T, p[x]$ ) Case 1
8.               $w \leftarrow \text{right}[p[x]]$  Case 1
9.          if  $\text{color}[\text{left}[w]] = \text{BLACK}$  and  $\text{color}[\text{right}[w]] = \text{BLACK}$ 
10.              $\text{color}[w] \leftarrow \text{RED}$  Case 2
11.              $x \leftarrow p[x]$  Case 2
12.          else if  $\text{color}[\text{right}[w]] = \text{BLACK}$ 
13.               $\text{color}[\text{left}[w]] \leftarrow \text{BLACK}$  Case 3
14.               $\text{color}[w] \leftarrow \text{RED}$  Case 3
15.              RIGHT-ROTATE( $T, w$ ) Case 3
16.               $w \leftarrow \text{right}[p[x]]$  Case 3
17.               $\text{color}[w] \leftarrow \text{color}[p[x]]$  Case 4
18.               $\text{color}[p[x]] \leftarrow \text{BLACK}$  Case 4
19.               $\text{color}[\text{right}[w]] \leftarrow \text{BLACK}$  Case 4
20.              LEFT-ROTATE( $T, p[x]$ ) Case 4
21.               $x \leftarrow \text{root}[T]$  Case 4
22.          else (same as then clause with “right” and “left” exchanged)
23.               $\text{color}[x] \leftarrow \text{BLACK}$ 
```

Case 1: w is red



- ▶ w must have black children.
- ▶ Make w black and $p[x]$ red.
- ▶ Then left rotate on $p[x]$.
- ▶ New sibling of x was a child of w before rotation \rightarrow must be black.
- ▶ Go immediately to case 2, 3, or 4.


Case 2: w is black & both of w 's children are black



- ▶ Take 1 black off x (\rightarrow singly black) and off w (\rightarrow red).
- ▶ Move that black to $p[x]$.
- ▶ Do the next iteration with $p[x]$ as the new x .
- ▶ If entered this case from case 1, then $p[x]$ was red \rightarrow new x is red & black \rightarrow color attribute of new x is RED \rightarrow loop terminates. Then new x is made black in the last line.


w is black,

Case 3: w 's left child is red, and w 's right child is black



- ▶ Make w red and w 's left child black.
- ▶ Then right rotate on w .
- ▶ New sibling w of x is black with a red right child → case 4.

Case 4: w is black, and w 's right child is red



- ▶ Make w be $p[x]$'s color (c).
- ▶ Make $p[x]$ black and w 's right child black.
- ▶ Then left rotate on $p[x]$.
- ▶ Remove the extra black on x (\Rightarrow x is now singly black) without violating any red-black properties.
- ▶ All done. Setting x to root causes the loop to terminate.

Time complexity of RB-DELETE

▶ **Analysis:**

- ▶ Case 2 is the only case in which more iterations occur.
 - ▶ x moves up 1 level.
 - ▶ Hence, $O(\lg n)$ iterations.
- ▶ Each of cases 1, 3, and 4 has 1 rotation $\rightarrow \leq 3$ rotations in all.
- ▶ Thus, the overall time for RB-DELETE is therefore $O(\lg n)$.
- ▶ <https://www.cs.usfca.edu/~galles/visualization/RedBlack.html>
(Red-black Tree Animation)